

Neural Network Hyperparameter Tuning for Flight Delay Classification

Pratik Ganguli

Title: Classification Analysis of Late Flight Arrivals

This example is adapted from <https://www.tidymodels.org/start/recipes/>

Load required packages

```
library(tidymodels)
```

```
-- Attaching packages ----- tidymodels 1.2.0 --
```

v broom	1.0.7	v recipes	1.1.0
v dials	1.3.0	v rsample	1.2.1
v dplyr	1.1.4	v tibble	3.2.1
v ggplot2	3.5.1	v tidyr	1.3.1
v infer	1.0.7	v tune	1.2.1
v modeldata	1.4.0	v workflows	1.1.4
v parsnip	1.2.1	v workflowsets	1.1.0
v purrr	1.0.2	v yardstick	1.3.2

```
-- Conflicts ----- tidymodels_conflicts() --
```

```
x purrr::discard() masks scales::discard()
```

```
x dplyr::filter() masks stats::filter()
```

```
x dplyr::lag() masks stats::lag()
```

```
x recipes::step() masks stats::step()
```

```
* Learn how to get started at https://www.tidymodels.org/start/
```

```
library(themis)
library(nycflights13)
#install.packages("nnet")
#install.packages("NeuralNetTools")
library(nnet)
library(NeuralNetTools)

tidymodels_prefer()
```

Question of interest: Will a plane departing from a New York City airport arrive more than 30 minutes late?

Data: all flights from New York City airports in 2013

Data Preparation

Data Preparation: Data Transformation 1

```
set.seed(123)
flight_data <-
  flights |>
  mutate(
    # Convert the arrival delay to a factor
    arr_delay = ifelse(arr_delay >= 30, "late", "on_time"),
    arr_delay = factor(arr_delay),
    # We will use the date (not date-time) in the recipe below
    date = lubridate::as_date(time_hour)
  ) |>
  # Include the weather data at origin
  inner_join(weather, by = c("origin", "time_hour")) |>
  # Take a random sample of flights (original dataset is too big)
  # Only do this for demo purposes.
  # On real data use the whole dataset
  slice_sample(n = 10000)
```

Data Transformation 2

```

flight_data <- flight_data |>
  # Only retain the specific columns we will use
  select(flight, time_hour, arr_delay,
         date, dep_time,
         #air_time, distance,
         #temp, carrier, origin, dest,
         ) |>
  # Exclude flights with missing arrival delay
  filter(!is.na(arr_delay)) |>
  # Encode qualitative columns as factors (instead of character strings)
  mutate_if(is.character, as.factor)

```

Note: We won't use variables `flight` and `time_hour` in the model, but keeping them in the dataset is useful for identification.

Modelling

Modelling: Pre-processing: Partition the data

80% training, 20% test

```

# Fix the random numbers by setting the seed
# This enables the analysis to be reproducible
# when random numbers are used
set.seed(222)

# Split 80% of the data into the training set
data_split <- initial_split(flight_data, prop = 0.8, strata = arr_delay)

# Create data frames for the three sets:
train_data <- training(data_split)
test_data <- testing(data_split)

```

Modelling: Pre-processing: Recipe

```

flights_rec <-
  # create recipe and specify formula
  recipe(arr_delay ~ ., data = train_data) |>
  # update role of ID variables

```

```

update_role(flight, time_hour, new_role = "ID") |>
# pre-process dates - extract day of week and month
step_date(date, features = c("dow", "month"),
          keep_original_cols = FALSE) |>
# normalize variables (required for knn)
step_normalize(all_numeric_predictors()) |>
# create dummy variables for nominal predictors
step_dummy(all_nominal_predictors())|>
# remove zero variance predictors
step_zv(all_predictors()) |>
# use upsampling to address class imbalance
step_upsample(arr_delay, over_ratio = 1)

```

Inspect the impact of the recipe

```

flights_prepped <-
flights_rec |>
  prep() |>
  bake(new_data = NULL)

flights_prepped |> glimpse()

```

```

Rows: 13,064
Columns: 21
$ flight      <int> 1109, 5699, 1532, 2454, 2189, 371, 471, 5378, 3069, 352~
$ time_hour   <dtm> 2013-02-07 17:00:00, 2013-05-27 14:00:00, 2013-08-15 2~
$ dep_time    <dbl> 0.831601469, 1.145951129, 1.986938534, 1.221476697, 1.7~
$ arr_delay   <fct> late, late, late, late, late, late, late, late, late, 1~
$ date_dow_Mon <dbl> 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0~
$ date_dow_Tue <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1~
$ date_dow_Wed <dbl> 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
$ date_dow_Thu <dbl> 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0~
$ date_dow_Fri <dbl> 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0~
$ date_dow_Sat <dbl> 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0~
$ date_month_Feb <dbl> 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0~
$ date_month_Mar <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
$ date_month_Apr <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0~
$ date_month_May <dbl> 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
$ date_month_Jun <dbl> 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0~
$ date_month_Jul <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0~
$ date_month_Aug <dbl> 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~

```

```
$ date_month_Sep <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
$ date_month_Oct <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0~
$ date_month_Nov <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0~
$ date_month_Dec <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1~
```

```
flights_prepped |> count(arr_delay)
```

```
# A tibble: 2 x 2
  arr_delay     n
  <fct>       <int>
1 late       6532
2 on_time    6532
```

Tuning: Specify model and workflow

Does the number of hidden nodes and number of iterations make a difference? `tidymodels` has some functions facilitate trying (tuning) different parameters.

Notes:

- The `mlp` has 1 hidden layer, so `hidden_units` specifies the number nodes in the hidden layer.
- epochs are the number of “passes” through the training data

Specify model

```
nn_model_tune <-
  mlp(hidden_units = tune(), # number of nodes in hidden layer
       epochs = tune() # number of iterations
  ) |>
  set_engine("nnet") |>
  set_mode("classification")
```

Create workflow = recipe + model

```
nn_tune_wf <- workflow() |>
  add_model(nn_model_tune) |>
  add_recipe(flights_rec)
```

Tuning: Define values to test

Create set of values to test

```
# grid_regular chooses sensible values for the parameters
nn_grid <- grid_regular(hidden_units(),
                        epochs(),
                        levels = c(5, 3))
print(nn_grid, n = 16)
```

```
# A tibble: 15 x 2
  hidden_units epochs
    <int>    <int>
1         1      10
2         3      10
3         5      10
4         7      10
5        10      10
6         1     505
7         3     505
8         5     505
9         7     505
10        10     505
11         1    1000
12         3    1000
13         5    1000
14         7    1000
15        10    1000
```

Tuning: Use cross validation to compare hyperparameter values

Create 5 cross-validation folds using training data

```
set.seed(747)
flights_folds <- vfold_cv(train_data, v = 5, strata = arr_delay)
flights_folds
```

```
# 5-fold cross-validation using stratification
# A tibble: 5 x 2
  splits          id
  <list>        <chr>
```

```

1 <split [6217/1555]> Fold1
2 <split [6217/1555]> Fold2
3 <split [6218/1554]> Fold3
4 <split [6218/1554]> Fold4
5 <split [6218/1554]> Fold5

```

Fit model on each fold

```

flights_metric <- metric_set(
  accuracy,
  roc_auc,
  sens, spec, bal_accuracy)

Sys.time()

```

```
[1] "2025-03-24 03:47:25 NZDT"
```

```

flights_res <- nn_tune_wkflow |>
  tune_grid(
    # select object containing folds
    resamples = flights_folds,
    # specify grid of values to evaluate
    grid = nn_grid,
    # specify set of metric (optional)
    metrics = flights_metric,
    # save predictions
    control = control_grid(save_pred = TRUE)
  )

Sys.time()

```

```
[1] "2025-03-24 03:48:36 NZDT"
```

Inspect results

```
flights_res
```

```

# Tuning results
# 5-fold cross-validation using stratification
# A tibble: 5 x 5
  splits          id    .metrics    .notes    .predictions

```

```

      <list>           <chr> <list>           <list>           <list>
1 <split [6217/1555]> Fold1 <tibble [75 x 6]> <tibble [0 x 3]> <tibble>
2 <split [6217/1555]> Fold2 <tibble [75 x 6]> <tibble [0 x 3]> <tibble>
3 <split [6218/1554]> Fold3 <tibble [75 x 6]> <tibble [0 x 3]> <tibble>
4 <split [6218/1554]> Fold4 <tibble [75 x 6]> <tibble [0 x 3]> <tibble>
5 <split [6218/1554]> Fold5 <tibble [75 x 6]> <tibble [0 x 3]> <tibble>

```

```

# inspect metrics for first fold
flights_res$.metrics[[1]]

```

```

# A tibble: 75 x 6

```

	hidden_units	epochs	.metric	.estimator	.estimate	.config
	<int>	<int>	<chr>	<chr>	<dbl>	<chr>
1	1	10	accuracy	binary	0.648	Preprocessor1_Model01
2	1	10	sens	binary	0.706	Preprocessor1_Model01
3	1	10	spec	binary	0.637	Preprocessor1_Model01
4	1	10	bal_accuracy	binary	0.671	Preprocessor1_Model01
5	1	10	roc_auc	binary	0.732	Preprocessor1_Model01
6	3	10	accuracy	binary	0.668	Preprocessor1_Model02
7	3	10	sens	binary	0.718	Preprocessor1_Model02
8	3	10	spec	binary	0.658	Preprocessor1_Model02
9	3	10	bal_accuracy	binary	0.688	Preprocessor1_Model02
10	3	10	roc_auc	binary	0.743	Preprocessor1_Model02

```

# i 65 more rows

```

Tuning: Evaluate performance of each k

Inspect results for each fold

```

flights_res |> collect_metrics(summarize = FALSE) |> print(n=7)

```

```

# A tibble: 375 x 7

```

	id	hidden_units	epochs	.metric	.estimator	.estimate	.config
	<chr>	<int>	<int>	<chr>	<chr>	<dbl>	<chr>
1	Fold1	1	10	accuracy	binary	0.648	Preprocessor1_Mod~
2	Fold1	1	10	sens	binary	0.706	Preprocessor1_Mod~
3	Fold1	1	10	spec	binary	0.637	Preprocessor1_Mod~
4	Fold1	1	10	bal_accuracy	binary	0.671	Preprocessor1_Mod~
5	Fold1	1	10	roc_auc	binary	0.732	Preprocessor1_Mod~
6	Fold2	1	10	accuracy	binary	0.719	Preprocessor1_Mod~
7	Fold2	1	10	sens	binary	0.560	Preprocessor1_Mod~

```

# i 368 more rows

```


Inspect summarised results

```
flights_res |> collect_metrics() |> print(n=7)
```

```
# A tibble: 75 x 8
```

	hidden_units	epochs	.metric	.estimator	mean	n	std_err	.config
	<int>	<int>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
1	1	10	accuracy	binary	0.669	5	0.0157	Preprocessor1~
2	1	10	bal_accuracy	binary	0.655	5	0.00610	Preprocessor1~
3	1	10	roc_auc	binary	0.707	5	0.00901	Preprocessor1~
4	1	10	sens	binary	0.636	5	0.0276	Preprocessor1~
5	1	10	spec	binary	0.675	5	0.0232	Preprocessor1~
6	3	10	accuracy	binary	0.686	5	0.0151	Preprocessor1~
7	3	10	bal_accuracy	binary	0.655	5	0.00968	Preprocessor1~

```
# i 68 more rows
```

Collect predictions

```
flights_pred <- flights_res |> collect_predictions()
flights_pred
```

```
# A tibble: 116,580 x 9
```

	.pred_class	.pred_late	.pred_on_time	id	.row	hidden_units	epochs
	<fct>	<dbl>	<dbl>	<chr>	<int>	<int>	<int>
1	on_time	0.474	0.526	Fold1	4	1	10
2	late	0.584	0.416	Fold1	6	1	10
3	on_time	0.464	0.536	Fold1	9	1	10
4	late	0.535	0.465	Fold1	12	1	10
5	late	0.564	0.436	Fold1	20	1	10
6	late	0.580	0.420	Fold1	23	1	10
7	late	0.571	0.429	Fold1	24	1	10
8	late	0.523	0.477	Fold1	32	1	10
9	late	0.584	0.416	Fold1	33	1	10
10	late	0.578	0.422	Fold1	34	1	10

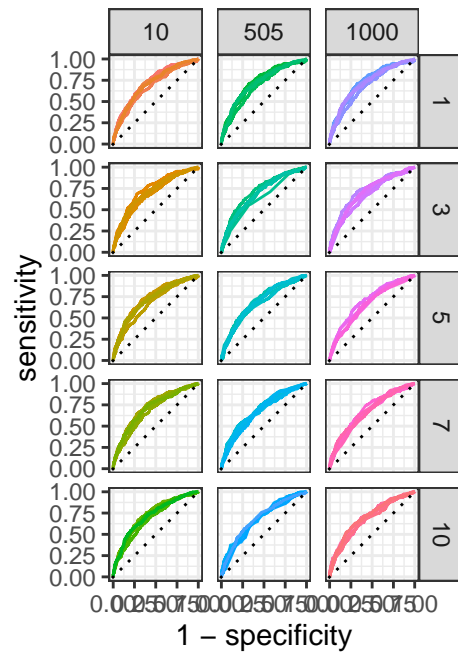
```
# i 116,570 more rows
# i 2 more variables: arr_delay <fct>, .config <chr>
```

Plot ROC curves

```

flights_pred |>
  group_by(id, hidden_units, epochs) |># id contains folds
  roc_curve(truth = arr_delay, .pred_late) |>
  autoplot() +
  facet_grid(rows = vars(hidden_units), cols = vars(epochs))+ theme(legend.position = "none")

```

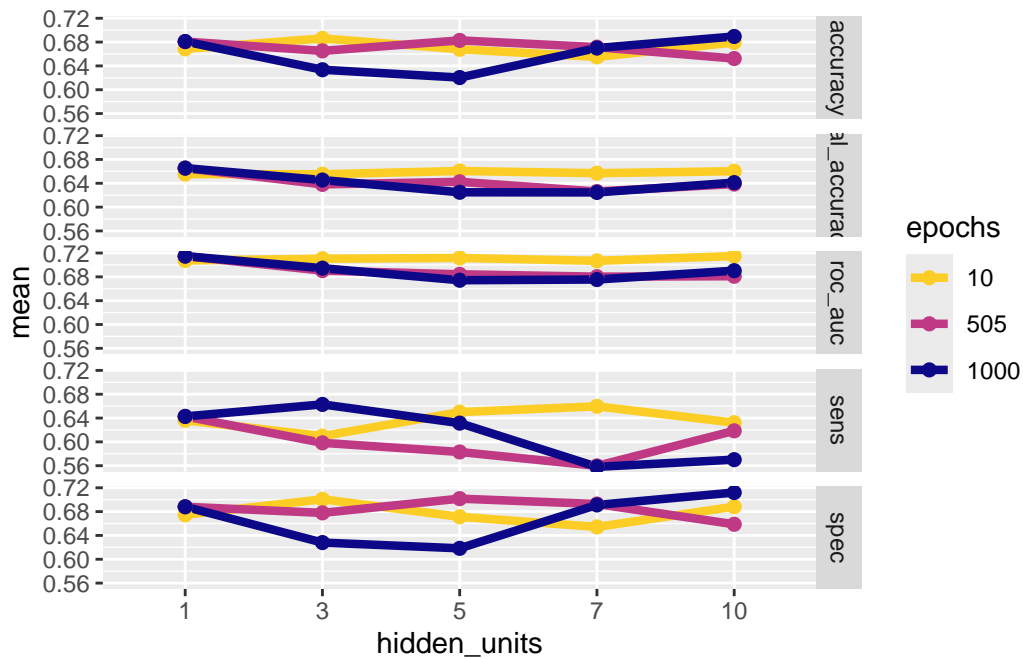


Plot metrics

```


g <-
flights_res |>
  collect_metrics() |>
  mutate(hidden_units = factor(hidden_units),
         epochs = factor(epochs)) |>
  ggplot(aes(hidden_units, mean, color = epochs, group = epochs)) +
  geom_line(linewidth = 1.5) +
  geom_point(size = 2) +
  scale_color_viridis_d(option = "plasma", begin = .9, end = 0)+
  facet_grid(rows = vars(.metric))
g

```



Alternative representation

```
g <-
flights_res |>
  collect_metrics() |>
  mutate(hidden_units = factor(hidden_units),
         epochs = factor(epochs)) |>
  ggplot(aes(x = .metric, y=mean, fill = .metric, group = .metric)) +
  geom_col(linewidth = 1.5) +
  #geom_point(size = 2) +
  scale_color_viridis_d(option = "plasma", begin = .9, end = 0)+
  facet_grid(rows = vars(hidden_units ), cols = vars(epochs))+ theme(axis.text.x = element_t
g
```



Tuning: Select the best model

In some applications, it is clear which metric should be used. In this case, use `select_best` with the specified metric.

However sometimes, it is worth making a trade-off to sacrifice the score on one metric, to avoid very poor score on another metric, or to reduce computational complexity (e.g. run-time or number of parameters). These types of trade-offs usually have to be performed manually. The following code demonstrates how to inspect the metrics.

```
flights_res |> select_best(metric = "roc_auc")
```

```
# A tibble: 1 x 3
  hidden_units epochs .config
    <int>    <int> <chr>
1         10     10 Preprocessor1_Model105
```

```
flights_res |> select_best(metric = "sens")
```

```
# A tibble: 1 x 3
  hidden_units epochs .config
```

```

      <int> <int> <chr>
1          3    1000 Preprocessor1_Model12

```

Inspect training performance for selected parameters

```

flights_res |>
  collect_metrics() |> filter(.config %in% c("Preprocessor1_Model02",
                                             "Preprocessor1_Model11",
                                             "Preprocessor1_Model12"))

```

A tibble: 15 x 8

	hidden_units	epochs	.metric	.estimator	mean	n	std_err	.config
	<int>	<int>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
1	3	10	accuracy	binary	0.686	5	0.0151	Preprocessor~
2	3	10	bal_accuracy	binary	0.655	5	0.00968	Preprocessor~
3	3	10	roc_auc	binary	0.710	5	0.0109	Preprocessor~
4	3	10	sens	binary	0.610	5	0.0352	Preprocessor~
5	3	10	spec	binary	0.701	5	0.0234	Preprocessor~
6	1	1000	accuracy	binary	0.681	5	0.00641	Preprocessor~
7	1	1000	bal_accuracy	binary	0.665	5	0.00999	Preprocessor~
8	1	1000	roc_auc	binary	0.715	5	0.0104	Preprocessor~
9	1	1000	sens	binary	0.643	5	0.0199	Preprocessor~
10	1	1000	spec	binary	0.688	5	0.00730	Preprocessor~
11	3	1000	accuracy	binary	0.634	5	0.0312	Preprocessor~
12	3	1000	bal_accuracy	binary	0.645	5	0.0135	Preprocessor~
13	3	1000	roc_auc	binary	0.695	5	0.0125	Preprocessor~
14	3	1000	sens	binary	0.663	5	0.0304	Preprocessor~
15	3	1000	spec	binary	0.628	5	0.0413	Preprocessor~

Select the best model

```

(selected_nn_model <-
  flights_res |>
    select_best(metric ="roc_auc")
    #select_best(metric ="sens")
)

```

A tibble: 1 x 3

	hidden_units	epochs	.config
	<int>	<int>	<chr>
1	10	10	Preprocessor1_Model105

```
## alternatively, manually specify selected model
# selected_nn_model <- tibble(
#   hidden_units = 5,
#   epochs = 10,
#   rowNumber = 2, #obtain from nn_grid
#   .config = paste0("Preprocessor1_Model", ifelse(rowNumber<10, "0", ""), rowNumber)
# )

flights_res |> collect_metrics()|> filter(.config == selected_nn_model$.config)
```

```
# A tibble: 5 x 8
  hidden_units epochs .metric      .estimator mean      n std_err .config
    <int>    <int> <chr>      <chr>      <dbl> <int>   <dbl> <chr>
1         10     10 accuracy    binary     0.679     5 0.0103 Preprocessor1~
2         10     10 bal_accuracy binary     0.660     5 0.00749 Preprocessor1~
3         10     10 roc_auc      binary     0.715     5 0.00946 Preprocessor1~
4         10     10 sens         binary     0.632     5 0.0183 Preprocessor1~
5         10     10 spec         binary     0.688     5 0.0139 Preprocessor1~
```

Tuning: Finalise the workflow

```
final_nn_tune_wkflow <-
  nn_tune_wkflow |>
  finalize_workflow(selected_nn_model)
```

Tuning: Final fit

Do a final fit (train on all training data and test on testing data)

```
final_fit <-
  final_nn_tune_wkflow |>
  last_fit(data_split,
    metrics = flights_metric)
```

Note: choice of metrics will vary for different applications.

Tuning: Evaluate final fit

```
final_fit |>  
  collect_metrics()
```

```
# A tibble: 5 x 4  
  .metric      .estimator .estimate .config  
  <chr>       <chr>      <dbl> <chr>  
1 accuracy    binary      0.701 Preprocessor1_Model1  
2 sens        binary      0.610 Preprocessor1_Model1  
3 spec        binary      0.719 Preprocessor1_Model1  
4 bal_accuracy binary      0.664 Preprocessor1_Model1  
5 roc_auc      binary      0.705 Preprocessor1_Model1
```

Tuning: Collect predictions and plot ROC curve

```
final_fit |>  
  collect_predictions() |>  
  roc_curve(truth = arr_delay, .pred_late) |>  
  autoplot()
```

